

# HarvestNet - Vegetable and Fruit Recognition using CNN - Group I

000

001

002

## A. Abstract

003

004

005

006

007

008

009

010

011

012

013

014

015

016

017

018

019

020

021

022

023

024

025

026

027

028

029

030

031

032

033

034

035

036

037

038

039

040

041

042

043

044

045

046

047

048

049

050

051

052

053

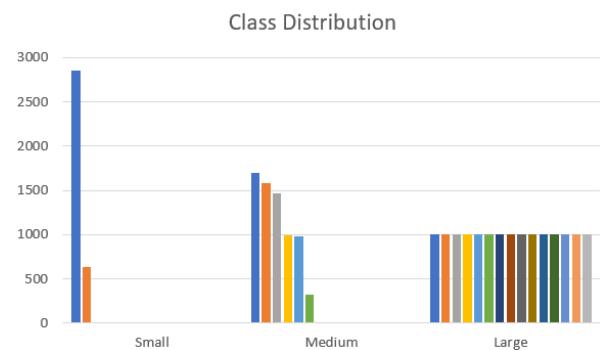
The aim of the project is to provide a comprehensive analysis on vegetable and fruits classification across three datasets that vary by number of classes on three models - ResNet18, VGGNet16 and AlexNet. Data transformation and augmentation is performed to remove class imbalance and improve intra-class diversity. Models are compared based on precision, recall, F1-score and accuracy. Two out of the nine models undergo transfer learning with ImageNet weights, leading to an increase in testing accuracy by 77 percent for VGGNet16 on large dataset. Further, hyperparameter tuning revealed that a lower learning rate of 10-4 and batch size of 128 is best suited for VGG16 and AlexNet. The performance of models with and without data augmentation is analysed, with the generation of T-SNE plots to understand dataset complexity. Overall, the performance of ResNet18 was the best across all three datasets. This report provides a detailed explanation for the observed results.

## B. Introduction

Food forms the basis of human existence. In recent years, the demand for food has increased drastically due to population growth, necessitating changes in agricultural methodologies. However, even now, a large number of tasks primarily related to produce picking, sorting and sales are done manually. This consumes a large amount of labour and is also very inefficient. Automated vegetable and fruit classification systems can help farmers in such a scenario. Additionally, classification can be used in inventory management sector to identify and restock produce correctly. It might also benefit kids and the blind as well as enhance grocery store self-checkouts. However, this classification problem is comparatively more challenging because of the produces' wide variety in terms of shape, color, and texture. In some cases, similar-looking produce may rely solely on texture for differentiation, necessitating a deep learning model capable of fine-grained classifications. Factors such as ripeness differences, can contribute to misclassification. Conversely, intra-class samples may contain too many similar images, potentially decreasing the generalisability of the model. One of the existing solutions to improve intra-class diversity is data augmentation in terms of rotations, scaling and color jittering. Further, transfer learning provides a standardised initial weights, resulting in model capable of identifying produce that look similar due to more fine-tuned weights. However, color jittering and excessive shape transformations may over distort the appearance of produce resulting in more misclassifications. Also, transfer learning may also lead to overfitting as the model may over fine-tune its weights. In order to classify these images effectively,

we have taken a systematic approach involving several key stages:

**1. Dataset Partitioning and models:** Initially, we gather three datasets of produce images that vary in number of classes, small, medium and large, ensuring a variability in angles. The datasets undergo transformations to enhance intra-class diversity and address class imbalance through data augmentation. CNN architectures such as, VGGNet16, AlexNet and ResNet18, evaluated through accuracy, precision, recall, and F1-score. Image samples are split into training, validation, and test sets in the ratio 80:10:10 for medium and small, while the large dataset was already split into train (70), validation (15) and test (15) folders.



**2. Initial Hyperparameter Exploration:** For the small and medium datasets, we initially experimented with a batch size of 64 and found that a learning rate of 0.01 with the Adam optimizer yielded the best results.

**3. Performance Analysis:** Upon testing these configurations on the large dataset, we observed varied performance across different CNN architectures. While ResNet performed well, VGG16 and AlexNet exhibited suboptimal results.

**4. Adaptive Strategy for Large Dataset:** In response to the performance disparity, we explored alternative optimization algorithms. Stochastic Gradient Descent (SGD), known for its momentum parameter, showed promise, particularly with VGG16 on the large dataset.

**5. Fair Comparison and Hyperparameter Tuning:** To maintain fairness in model comparison, we retained the Adam optimizer with a learning rate of 0.01 for the large dataset across all models. However, extensive hyperparameter tuning was conducted, including experimentation with batch sizes and learning rates, detailed in the ablative study within the results section. Various metrics were considered to evaluate the algorithm's performance:

**Accuracy:** The proportion of correctly classified samples.

**Precision:** The ratio of correctly predicted positive observations to the total predicted positives.

**Recall:** The ratio of correctly predicted positive observations to the actual positives.

054

055

056

057

058

059

060

061

062

063

064

065

066

067

068

069

070

071

072

073

074

075

076

077

078

079

080

081

082

083

084

085

086

087

088

089

090

091

092

093

094

095

096

097

098

099

100

101

102

103

104

105

106

107

108     **F1 Score:** The harmonic mean of precision and recall, providing a balance between the two.  
109     **Mean Loss:** The average loss computed during training or evaluation. Models trained from scratch are expected to have much lower accuracy and higher loss than models trained with transfer learning

## 115     B.1. Related Works

117     In this research [13], they improved the performance of vanilla VGGNet model on vegetable classification by combining the outputs of two fully connected layers that achieved an accuracy of 95.8 percent. They further improved this model by including batch normalisation layers that achieved 96.5 percent accuracy. Study in [14] uses image saliency to detect most important features. They used VGGNet and achieved an accuracy of 95.6 percent. Researchers in [15] used a deep neural network that used stochastic gradient descent that processed 3 million learning iterations to achieve an accuracy of 97.58 percent, however, more than 10 million learning iterations resulted in performance decline. Work in [16] entailed vegetable classification using CNN but focuses on reducing misclassifications when vegetables look similar in terms of colour. In paper [17], used a variety of CNN based models pretrained vs trained from scratch on VGGNet16, MobileNet, InceptionV3 and ResNet to ascertain the best performing model. Overall, they concluded that transfer learning produced the best model with over 99.9 percent accuracy in InceptionV3 and ResNet50. In this paper [18], they used a modified AlexNet model that used ReLU instead of the sigmoid function after the output layer that sped up the training process. They achieved a testing accuracy of 92.1 percent.

## 141     C. Methodology

### 142     C.1. Datasets

145     We have used 3 different datasets from kaggle for our study:

- 147     1. Vegetable Image Dataset [Large dataset]
- 148     2. Fresh and Stale Images of Fruits and Vegetables [Medium dataset]
- 149     3. Fresh and Rotten Classification [Small dataset]

152     The large dataset consists of 21,000 images from 15 different classes. Each class of vegetable has a uniform distribution of 1400 images. The dataset is split into 70 percent for training (15,000 images), 15 percent for validation (3,000) and 15 percent for testing (3,000) purposes. The image resolution for the images in the dataset is 224x224. The original medium dataset contains 14 classes and is down-sized to 6. There are a total of 7038 images among 6 classes. The small dataset contains 14 classes but we reduced it to 2 classes. There are a total of 3486 images

162     among 2 chosen classes. The dataset is divided into training (2787 samples), testing (348 samples) and validation set (349 samples) in the ratio 80:10:10. All three datasets have varied class representation with rotated images, different illuminations, cut and uncut produce, produce found in different background setting such as market place, kitchen, cutting board, basket and a person holding it. There are also variation in illuminations and the number of produce in a single image for example, some images have a single vegetable whereas other images have at least 7-8 vegetables stacked together. Further, different variants of a produce such as apple (green and red apples) and tomato are included in the dataset to improve the robustness of the trained CNN model.



(a) Bean



(b) Brinjal



(c) Capsicum

Figure 1. Vegetables of Large Dataset



(a) fresh apple



(b) fresh bitter gourd

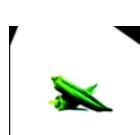


(c) fresh capsicum

Figure 2. Vegetables of Medium Dataset



(a) fresh banana



(b) freshokra

Figure 3. Vegetables of Small Dataset

195     All the three datasets are preprocessed using PyTorch transforms, resizing them to 224x224, flipping horizontally and vertically, and rotating by 10 degrees. Random erasing and colour jittering are applied to introduce variability in colour distribution, aiding the model's ability to recognize vegetables under various lighting conditions. These transformations are combined using ‘transforms.compose()’ to enhance dataset diversity and mitigate overfitting. Additionally, the medium and small datasets, due to class imbalance, undergo data augmentation using PyTorch’s WeightedRandomSampler. First, class weights are calculated inversely proportional to the number of samples in each class.

196  
197  
198  
199  
200  
201  
202  
203  
204  
205  
206  
207  
208  
209  
210  
211  
212  
213  
214  
215

These weights are then converted to a Double Tensor and passed to WeightedRandomSampler, which returns a class sampler. This sampler facilitates the sampling of data points during training, ensuring that the probability of selecting each sample is proportional to its assigned weight enabling a more balanced representation of classes during the training process. The sampler and training set is then passed as an argument to the DataLoader class to initialise the data loaders for the training set. The Testing set and Validation set do not undergo class balancing. The batch size chosen across all three datasets and three models is 64 initially.

## C.2. CNN Models

We have chosen ResNet-18, VggNet-16 and AlexNet for the vegetable classification with their standard layers. Other models like Inception-v3 and DenseNet have high computational complexity compared to ResNet and Vgg which could raise the problem of overfitting on the vegetable dataset. Lighter models like MobileNet might not generalise the images accurately given its trade off between accuracy and computational efficiency. Therefore, considering the complexity of our datasets, the three chosen models seem fit to the vegetable classification task.

**1. ResNet18:** ResNet18 has 18 layers, with a 7x7 kernel in the first layer, followed by four sets of Convolutional layers with two residual blocks. Every block is composed of two weight layers, and a skip connection that links the output of the second weight layer to the input using a ReLU activation function. The model accepts inputs of size 224x224. A deep model are more likely to capture complex features resulting in a more robust and fine-tuned model. This is imperative to our domain as produce images have high inter-class similarity in terms of color and shape. Skip connections also mitigates vanishing gradient problem, thus resulting in higher performance with an increase in model complexity or number of hidden layers. Additionally, despite ResNet18 having fewer parameters, it results in good performance and better training speeds due to residual connections. ResNet18 has 1.8 Billion flop calculations.

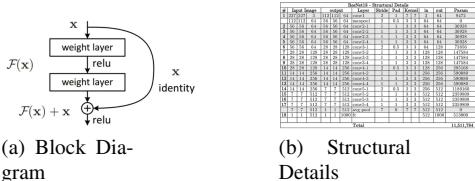


Figure 4. ResNet-18

**2. VGGNet16:** VGGNet16 has 16 layers and has 13 convolutional layers, 5 max pooling layers, and 2 dense layers. It operates on input of size 224x224 with 3 RGB

channels. The convolutional layers use a 3x3 kernel with a stride of 1 and padding of 1 and max-pooling layers have a 2x2 kernel with a stride of 2. The model has similar depth to ResNet18 but focuses on simplicity and does not have skip connections. We selected VggNet to assess whether simplicity in architecture yields comparable results. VGGNET16 has 15.3 billion flop calculations.

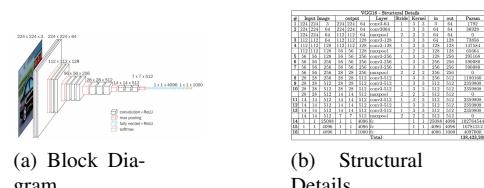


Figure 5. VGGNet-16

**3. AlexNet:** AlexNet is a light weight model with just eight layers, made up of convolutional layers, max-pooling layers and 3 fully connected layers. It uses ReLU activation function that enables faster convergence of minimum loss during model training. The model also has local response normalisation layers that enhances its generalisability by normalising the outputs across adjacent layers. We opted for AlexNet to ascertain whether a less complex model is suitable for our dataset. AlexNet has 714.2 million flop calculations.

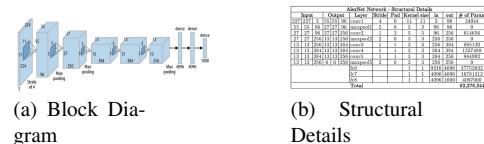


Figure 6. AlexNet

To analyse computation complexities, all three CNN models were trained on three datasets. Trends on medium, large and small datasets were similar.



Figure 7. Training (left) and Validation (right) accuracy for one-epoch training

This suggests that the ResNet18's initial weights and architecture are better suited for the problem statement. This can be attributed to deeper layers of ResNet18 and overall, more complex architecture than AlexNet and VggNet16. However, this does not imply that the model will produce the best validation accuracy at the end of model training

phase, but in our case ResNet18 did produce the best results compared to the other two models. From table 1, ResNet18 required the most time for one epoch training for all datasets except medium though less learnable parameters for a model that is 18 layers deep. This may be due to skip connections requiring more resources and computation that has lead to an increase in training time.

Table 1. Wall Clock time for one-epoch training

Dataset	ResNet18	VGGNet16	AlexNet
Small	152.09s	55.094s	45.104s
Medium	65.40s	166.82s	144.160s
Large	1881.07s	65.67s	36.98s

### C.3. Optimization Algorithm

To optimize the CNN models for the vegetable and fruit classification task, we employed a systematic approach that involved both validation and optimization. All CNN architectures were evaluated using metrics such as accuracy, precision, recall, and F1-score. Through this validation and optimization process, we aimed to identify the most effective optimization algorithm and hyperparameter configurations for each dataset size and CNN architecture combination, ensuring robust performance across various scenarios. Results are documented in section D.

For the small and medium datasets, as well as for consistency in comparison, the Adam optimizer with a learning rate of 0.01 was selected. Adam is preferred for its adaptive learning rate capabilities, which adjust the learning rates individually for each parameter based on their gradients' first and second moments. We have also experimented with Stochastic gradient descent (SGD) for VGG.

### Properties and Hyperparameters:

Optimiser	Adam	SGD
Properties	Adaptive learning rate, momentum, bias correction	Learning rate (lr), momentum, weight decay
Hyperparameters	Learning rate (lr), beta parameters (betas), epsilon (eps)	Learning rate (lr), beta parameters (betas), epsilon (eps)

## D. Results

### D.1. Experiment Setup

Various techniques were used to improve the vegetable and fruits image classification problem, including resizing

images, normalizing pixel values, and enhancing data through rotations, flips, and brightness adjustments. Random weighted sampling was found to be the most effective for most models. However, a batch size of 128 resulted in poorer test accuracy for AlexNet and ResNet18 for medium and small datasets. The small dataset showed a validation accuracy of 99.43 percent, while the medium dataset showed a validation accuracy of 92.89 percent. VggNet16 did not complete 10 epochs, resulting in a low test accuracy of 22.19 percent. For Large Dataset, the initial training was done using the same parameters as was used in small and medium dataset for a fair comparison.

### D.2. Main Results

**Small dataset:** AlexNet and VGGNet16 have a validation and testing accuracy of 18.34 percent, 18.97 percent and 76.85 percent 80.97 percent, respectively. This accuracy dramatically increases to 99.85 percent and 98.43 percent for the validation and testing accuracy for the ResNet respectively could be due to the skip connections that handles the vanishing gradient in an effective manner. ResNet's 18 layers, more than VGGNet and AlexNet, is one of the other reasons for higher accuracy value as the image samples are complex in nature. We observe that the validation accuracy over 10 epochs for the VGGNet and AlexNet alternate between high and low values. The probable reason is that the learning rate is higher than the optimal learning rate necessary to find the local minima, thus missing the local minima. The other reason could be that batch normalisation is not performed, which explains the fluctuation in the accuracies and leading to very large convergence time.

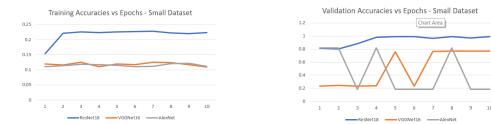


Figure 8. Training (on the left) and Validation (on the right) accuracy vs Epoch on Small Dataset

**Medium dataset:** ResNet18 has a validation accuracy of 92.6 percent. Whereas, AlexNet and VggNet16 produced a low validation accuracies.

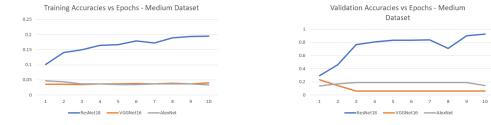


Figure 9. Training (on the left) and Validation (on the right) accuracy vs Epoch on Medium Dataset

This suggests that ResNet18 is able to capture the

most important features in the training samples better than AlexNet and VggNet16. In contrast, VGGNet does not have skip connections, making it more prone to vanishing gradient problems. AlexNet, on the other hand, is a light weight model which may be insufficient to capture the complexity of the training data as there is high intra-class and inter-class similarity.

**Large dataset:** The validation and the testing accuracies for the VGGNet and AlexNet is low at 6.67 percent for both. ResNet performs the best with the high score of 97.53 percent for the validation accuracy. The low performance of AlexNet and VGGNet stems from the problem of vanishing gradient which is mitigated in the case of ResNet18. The hyperparameters set for the models may not have worked well to accommodate large image sets.



Figure 10. Training (on the left) and Validation (on the right) accuracy vs Epoch on Large Dataset

Table 2. Comparison of 9 models trained from scratch

Sl. No.	Dataset	Model	Accuracy
1	Small	AlexNet	18.97
2	Small	VGG-16	80.97
3	Small	ResNet-18	99.43
4	Medium	AlexNet	15.22
5	Medium	VGG-16	4.5
6	Medium	ResNet-18	92.89
7	Large	AlexNet	6.67
8	Large	VGG-16	6.67
9	Large	ResNet-18	92.49

**Transfer Learning:** We then proceeded with transfer learning on the large dataset using VGGNet16 and AlexNet as they gave lower accuracy. The model run with VGGNet witnesses a dramatic increase in the test and the validation accuracies with the score of 95.07 percent and 95.6 percent. The likely reason for better results is that the model has better chance of tuning the weights when the weights have already been initialized with the transfer learning. We observe similar results with the AlexNet model, when the learning rate is reduced to 0.001. This tuning shows the increase of the test and the validation accuracies to 84.23 percent and 84.2 percent. We experimented with different learning rates

to see the impact, and found that a lower learning rate was better here.

Table 3. Comparison for Transfer Learning - Large Dataset

CNN Model	Transfer Learning	Learning Rate	Optimiser	Accuracy
VGGNet	True	0.01	Adam	85.97
AlexNet	True	0.01	Adam	6.67
AlexNet	False	0.001	Adam	76.9
AlexNet	True	0.001	Adam	84.23
VGGNet	False	0.001	SGD	79
VGGNet	True	0.001	SGD	99

**Performance of model with class imbalance:** We utilised the ResNet18 model on medium dataset to study the impact of data augmentation that uses weighted random sampler to remove class imbalance. ResNet18 produced a testing accuracy of 92.8 percent and 90.75 percent with and without data augmentation respectively. Thus, we can infer that there is a small performance improvement as class weights are inversely proportional to their sample size, which allows the model to learn from the minority classes as effectively as from the majority classes. The recall for minority classes such as fresh bitter gourd and fresh capsicum are only 37.5 percent and 80.3 percent for the imbalanced dataset and 100 percent and 91 percent for balanced dataset.

**T-SNE Plots:** TSNE plots on medium dataset are plotted after passing image data through the encoder (convolution) layers.

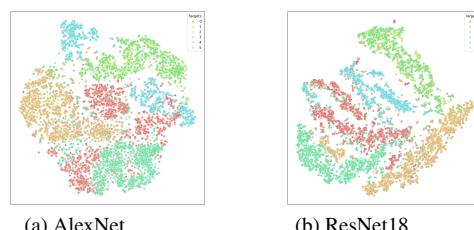


Figure 11. TSNE on Medium Dataset

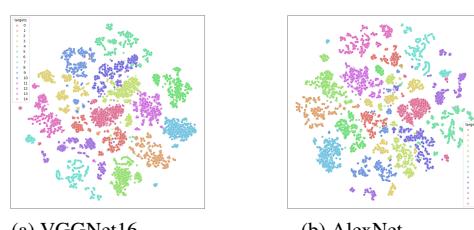


Figure 12. TSNE with Transfer Learning on Large Dataset

We can observe that the class clusters are more separable and distinct in ResNet18 than in AlexNet. This can be attributed to AlexNet model being less complex than ResNet18 and hence can not extract the high level features that distinguish the different types of produce in the dataset. On the other hand, VGGNet16 and AlexNet models that used pre-trained weights on Large Dataset produce very distinct clusters, better than models that have weights trained from scratch. This is because, transfer learning provides a better initial feature weights. Overall, ResNet18 performs the best in all the 3 cases. However, we observed very low accuracy using AlexNet and VGGNet16 especially for the large dataset. We applied transfer learning with these two models and achieved much better accuracy. For increasing the performance without transfer learning, we experimented with different learning rates and optimisers.

### D.3. Ablative Study

Along with just training the 9 models with fixed batch size of 64, learning rate = 0.01, plus 2 transfer learning models, we also performed hyperparameter tuning, particularly using Large Dataset for VGGNet using different learning rates and optimizers. While Adam performs well in terms of optimization, it can lead to worse generalization performance (i.e., poorer performance on unseen data) compared to SGD in image classification tasks when the optimal learning rate and batch size is not provided. With our selected batch size of 64, and 0.01 learning rate, SGD performed better. As SGD performs better than Adam in this context, we can infer from the t-SNE plot that there are tighter or more distinct clustering of data points corresponding to SGD compared to those corresponding to Adam. Tighter clusters suggest that SGD has resulted in features that better separate different vegetable classes in the reduced-dimensional space. Adam's performance might be affected by its inner hyperparameters, which can be computationally expensive to tune.

Table 4. VGG16 - Large Dataset - SGD vs Adam

Optimiser	Learning Rate	Transfer Learning	Accuracy
Adam	0.01	False	6.67
SGD	0.01	False	83.83

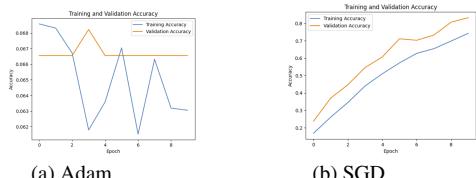


Figure 13. Accuracy Results on Large Dataset

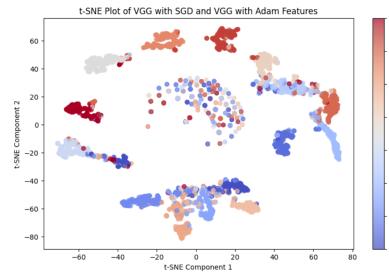


Figure 14. TSNE Results - Adam vs SGD

We also trained AlexNet using different batch sizes and learning rates, as depicted in the graph. We observed that for the large dataset, a batch size of 128 performs the best, while 8 performs the worst (using a fixed Learning rate of 0.01) in Alexnet, and a learning rate of 0.0001 gives best results as opposed to 0.1 which gives the worst accuracy (with fixed Batch Size of 64). For VGGNet also, similar pattern is observed for 5 different learning rates.

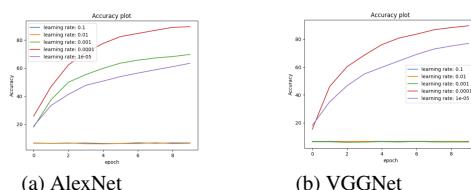


Figure 15. Accuracy Plot for different learning rates

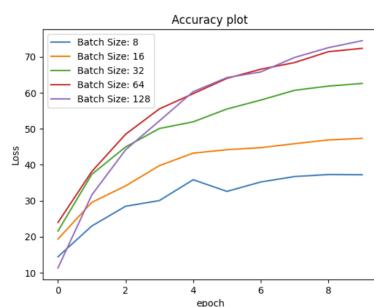


Figure 16. Large dataset - AlexNet - different batch sizes

This is in line with our expectations as larger batch sizes tend to provide more stable gradients during training by reducing the variance in gradient estimates, leading to more consistent weight updates. In case of learning rates, a high learning rate causes large weight updates, which can destabilize training while a very low learning rate leads to small weight updates, slowing down convergence. Hence, tuning is essential, as it helps us identify the best values of the different hyperparameters and see what works best for our model and dataset.

540  
541  
542  
543  
544  
545  
546  
547  
548  
549  
550  
551  
552  
553  
554  
555  
556  
557  
558  
559  
560  
561  
562  
563  
564  
565  
566  
567  
568  
569  
570  
571  
572  
573  
574  
575  
576  
577  
578  
579  
580  
581  
582  
583  
584  
585  
586  
587  
588  
589  
590  
591  
592  
593  
594  
595  
596  
597  
598  
599  
600  
601  
602  
603  
604  
605  
606  
607  
608  
609  
610  
611  
612  
613  
614  
615  
616  
617  
618  
619  
620  
621  
622  
623  
624  
625  
626  
627  
628  
629  
630  
631  
632  
633  
634  
635  
636  
637  
638  
639  
640  
641  
642  
643  
644  
645  
646  
647

648  
649  
650  
651  
652  
653  
654  
655  
656  
657  
658  
659  
660  
661  
662  
663  
664  
665  
666  
667  
668  
669  
670  
671  
672  
673  
674  
675  
676  
677  
678  
679  
680  
681  
682  
683  
684  
685  
686  
687  
688  
689  
690  
691  
692  
693  
694  
695  
696  
697  
698  
699  
700  
701

## References

- [1] Ahmed, M. I., Mamun, S. M., Asif, A. U. Z. (2021, May). Dcnn-based vegetable image classification using transfer learning: A comparative study. In 2021 5th International Conference on Computer, Communication and Signal Processing (ICCCSP) (pp. 235-243). IEEE.
- [2] Tapia-Mendez, E., Cruz-Albaran, I. A., Tovar-Arriaga, S., Morales-Hernandez, L. A. (2023). Deep Learning-Based Method for Classification and Ripeness Assessment of Fruits and Vegetables. *Applied Sciences*, 13(22), 12504.
- [3] Li, Z., Li, F., Zhu, L., Yue, J. (2020). Vegetable recognition and classification based on improved VGG deep learning network model. *International Journal of Computational Intelligence Systems*, 13(1), 559-564.
- [4] Hameed, K., Chai, D., Rassau, A. (2018). A comprehensive review of fruit and vegetable classification techniques. *Image and Vision Computing*, 80, 24-44.
- [5] Singh, H., Singh, R., Goel, P., Singh, A., Sharma, N. (2022). Automatic Framework for Vegetable Classification using Transfer-Learning. *Int. J. Electr. Electron. Res.*, 10(2), 405-410.
- [6] Vegetable Image Dataset: <https://www.kaggle.com/datasets/misrakahmed/vegetable-image-dataset>
- [7] Fresh and Stale Images of Fruits and Vegetables Dataset: <https://www.kaggle.com/datasets/raghavrpotdar/fresh-and-stale-images-of-fruits-and-vegetables>
- [8] Fresh and Rotten Classification Dataset <https://www.kaggle.com/datasets/swoyam2609/fresh-and-stale-classification>
- [9] CNN and Transfer learning: Vegetable Image Classification <https://medium.com/@rodrigonev98/cnn-and-transfer-learning-vegetable-image-classification-ecdc1dalfff>
- [10] Base on Resnet <https://www.kaggle.com/code/yuchuanwang/base-on-resnet-99-93-accuracy>
- [11] Vegetable Classifier - VGG <https://www.kaggle.com/code/ashwinshetgaonkar/vegetable-classifier-vgg-pytorch>
- [12] Difference between AlexNet, VGGNet, ResNet, and Inception <https://towardsdatascience.com/the-w3h-of-alexnet-vggnet-resnet-and-inception-7baaaecc96>
- [13] Li, Z., Li, F., Zhu, L., Yue, J. (2020). Vegetable recognition and classification based on improved VGG deep learning network model. *International Journal of Computational Intelligence Systems*, 13(1), 559-564.
- [14] Zeng, G. (2017, October). Fruit and vegetables classification system using image saliency and convolutional neural network. In 2017 IEEE 3rd Information Technology and Mechatronics Engineering Conference (ITOEC) (pp. 613-617). IEEE.
- [15] Sakai, Y., Oda, T., Ikeda, M., Barolli, L. (2016, July). A vegetable category recognition system using deep neural network. In 2016 10th International Conference on Innovative Mobile and Internet Services in Ubiquitous Computing (IMIS) (pp. 189-192). IEEE.
- [16] Duth, P. S., Jayasimha, K. (2020, May). Intra class vegetable recognition system using deep learning. In 2020 4th International conference on intelligent computing and control systems (ICICCS) (pp. 602-606). IEEE.
- [17] Ahmed, M. I., Mamun, S. M., Asif, A. U. Z. (2021, May). DCNN-based vegetable image classification using transfer learning: A comparative study. In 2021 5th International Conference on Computer, Communication and Signal Processing (ICCCSP) (pp. 235-243). IEEE.
- [18] Zhu, L., Li, Z., Li, C., Wu, J., Yue, J. (2018). High performance vegetable classification from images based on alexnet deep learning model. *International Journal of Agricultural and Biological Engineering*, 11(4), 217-223.
- [19] Resnet18 Model With Sequential Layer For Computing Accuracy On Image Classification Dataset <https://ijcrt.org/papers/IJCRT2205235.pdf>
- [20] Everything you need to know about VGG16 <https://medium.com/@mygreatlearning/everything-you-need-to-know-about-vgg16-7315defb5918>
- 702  
703  
704  
705  
706  
707  
708  
709  
710  
711  
712  
713  
714  
715  
716  
717  
718  
719  
720  
721  
722  
723  
724  
725  
726  
727  
728  
729  
730  
731  
732  
733  
734  
735  
736  
737  
738  
739  
740  
741  
742  
743  
744  
745  
746  
747  
748  
749  
750  
751  
752  
753  
754  
755

756

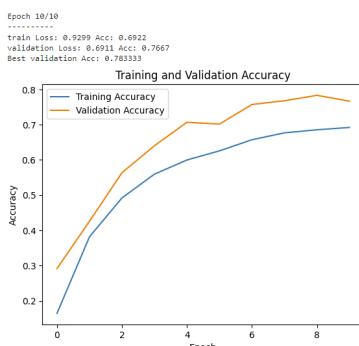
**E. Supplementary Material**

Figure 17. AlexNet - No transfer learning - 0.001 Learning Rate, Batch size 64 - LargeDataset - Train vs Validation accuracy

757

758

759

760

761

762

763

764

765

766

767

768

769

770

771

772

773

774

775

776

777

778

779

780

781

782

783

784

785

786

787

788

789

790

791

792

793

794

795

796

797

798

799

800

801

802

803

804

805

806

807

808

809

810

811

812

813

814

815

816

817

818

819

820

821

822

823

824

825

826

827

828

829

830

831

832

833

834

835

836

837

838

839

840

841

842

843

844

845

846

847

848

849

850

851

852

853

854

855

856

857

858

859

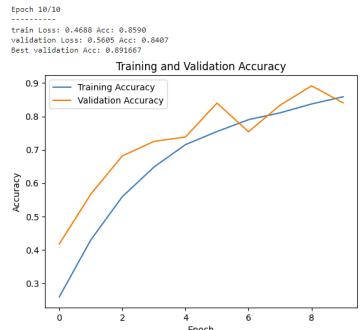


Figure 20. ResNet - No transfer learning - 0.01 Rate, Batch size 64 - LargeDataset - Train vs Validation accuracy

	precision	recall	f1-score	support
Bean	0.50	0.56	0.53	200
Bitter_Gourd	0.69	0.78	0.73	200
Bottle_Gourd	0.96	0.63	0.76	200
Brinjal	0.70	0.69	0.70	200
Broccoli	0.75	0.82	0.78	200
Cabbage	0.76	0.61	0.68	200
Capsicum	0.98	0.81	0.89	200
Carrot	0.94	0.94	0.94	200
Cauliflower	0.89	0.77	0.82	200
Cucumber	0.64	0.74	0.69	200
Papaya	0.78	0.83	0.80	200
Potato	0.79	0.99	0.88	200
Pumpkin	0.66	0.91	0.76	200
Radish	0.85	0.85	0.85	200
Tomato	0.92	0.61	0.74	200
accuracy			0.77	3000
macro avg	0.79	0.77	0.77	3000
weighted avg	0.79	0.77	0.77	3000

Figure 18. AlexNet - No transfer learning - 0.001 Learning Rate, Batch size 64 - LargeDataset - Metrics

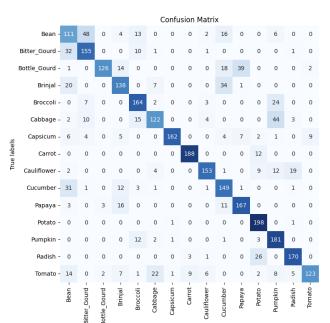


Figure 19. AlexNet - No transfer learning - 0.001 Learning Rate, Batch size 64 - LargeDataset - Confusion Matrix

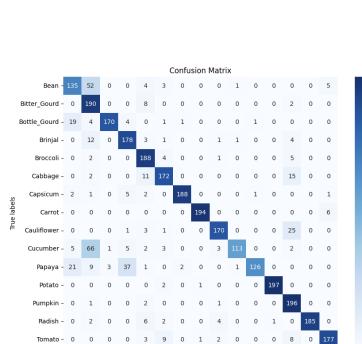


Figure 22. ResNet - No transfer learning - 0.01 Rate, Batch size 64 - LargeDataset - Confusion Matrix

## COMP 6721 Winter 2024 Submission #Group I. CONFIDENTIAL REVIEW COPY. DO NOT DISTRIBUTE.

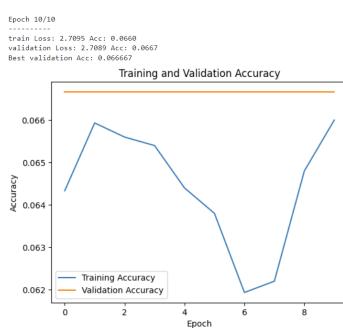


Figure 23. AlexNet - No transfer learning - 0.01 Rate, Batch size 64 - LargeDataset - Train vs Validation accuracy

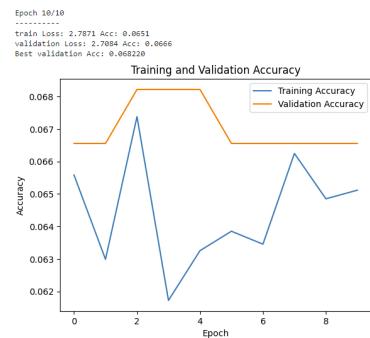


Figure 26. VGGNet - No transfer learning - 0.01 Rate, Batch size 64 - LargeDataset - Train vs Validation accuracy

	precision	recall	f1-score	support
Bean	0.00	0.00	0.00	200
Bitter_Gourd	0.00	0.00	0.00	200
Bottle_Gourd	0.00	0.00	0.00	200
Brinjal	0.00	0.00	0.00	200
Broccoli	0.00	0.00	0.00	200
Cabbage	0.00	0.00	0.00	200
Capiscum	0.00	0.00	0.00	200
Carrot	0.00	0.00	0.00	200
Cauliflower	0.00	0.00	0.00	200
Cucumber	0.00	0.00	0.00	200
Papaya	0.00	0.00	0.00	200
Potato	0.00	0.00	0.00	200
Pumpkin	0.00	0.00	0.00	200
Radish	0.00	0.00	0.00	200
Tomato	0.07	1.00	0.12	200
accuracy			0.07	3000
macro avg	0.00	0.07	0.01	3000
weighted avg	0.00	0.07	0.01	3000

Figure 24. AlexNet - No transfer learning - 0.01 Rate, Batch size 64 - LargeDataset - Metrics

	precision	recall	f1-score	support
Bean	0.00	0.00	0.00	200
Bitter_Gourd	0.00	0.00	0.00	200
Bottle_Gourd	0.00	0.00	0.00	200
Brinjal	0.00	0.00	0.00	200
Broccoli	0.00	0.00	0.00	200
Cabbage	0.00	0.00	0.00	200
Capiscum	0.00	0.00	0.00	200
Carrot	0.00	0.00	0.00	200
Cauliflower	0.00	0.00	0.00	200
Cucumber	0.00	0.00	0.00	200
Papaya	0.00	0.00	0.00	200
Potato	0.07	1.00	0.12	200
Pumpkin	0.00	0.00	0.00	200
Radish	0.00	0.00	0.00	200
Tomato	0.00	0.00	0.00	200
accuracy			0.07	3000
macro avg	0.00	0.07	0.01	3000
weighted avg	0.00	0.07	0.01	3000

Figure 27. VGGNet - No transfer learning - 0.01 Rate, Batch size 64 - LargeDataset - Metrics

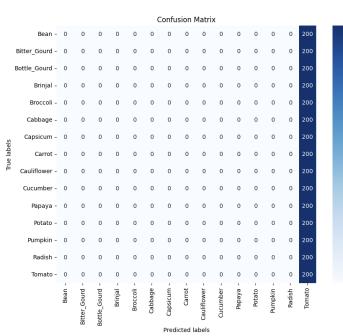


Figure 25. AlexNet - No transfer learning - 0.01 Rate, Batch size 64 - LargeDataset - Confusion Matrix

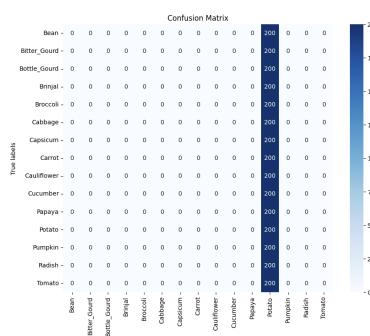


Figure 28. VGGNet - No transfer learning - 0.01 Rate, Batch size 64 - LargeDataset - Confusion Matrix

## COMP 6721 Winter 2024 Submission #Group I. CONFIDENTIAL REVIEW COPY. DO NOT DISTRIBUTE.

972

973

974

975

976

977

978

979

980

981

982

983

984

985

986

987

988

989

990

991

992

993

994

995

996

997

998

999

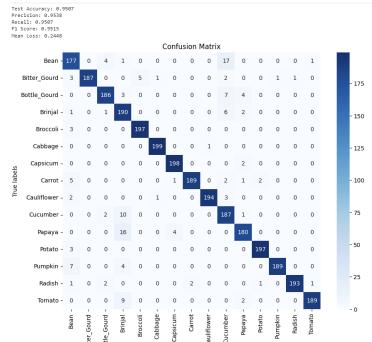
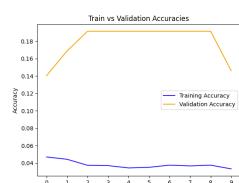


Figure 29. VGGNet - WithTransferLearning - Batch64 - LR0.01 - LargeDataset - Metrics and Confusion Matrix



```
from sklearn.metrics import classification_report
print(classification_report(true_labels, predicted_labels, target_names=['fresh_apple', 'fresh_banana', 'fresh_bitter_gourd',
    'fresh_capsicum', 'fresh_carrot', 'fresh_cauliflower', 'fresh_cucumber', 'fresh_papaya', 'fresh_potato',
    'fresh_pumpkin', 'fresh_radish', 'fresh_tomato'],
    digits=4))

accuracy: 0.0354
precision: 0.0354
recall: 0.0354
f1-score: 0.0354
mean_loss: 0.5686
```

Figure 30. AlexNet - NoTransferLearning - Batch64 - LR0.01 - Medium - Metrics and Accuracy curve

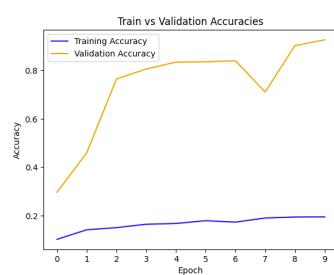


Figure 31. ResNet - NoTransferLearning - Batch64 - LR0.01 - Medium - Train vs Validation accuracy

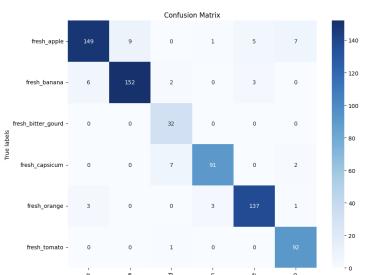


Figure 32. ResNet - NoTransferLearning - Batch64 - LR0.01 - Medium - Metrics

	precision	recall	f1-score	support
fresh_apple	0.9430	0.8713	0.9058	171
fresh_banana	0.9441	0.9325	0.9383	163
fresh_bitter_gourd	0.7619	1.0000	0.8649	32
fresh_capsicum	0.9579	0.9100	0.9333	100
fresh_orange	0.9448	0.9514	0.9481	144
fresh_tomato	0.9020	0.9892	0.9436	93
accuracy			0.9289	703
macro avg	0.9090	0.9424	0.9223	703
weighted avg	0.9321	0.9289	0.9290	703

Figure 33. ResNet - NoTransferLearning - Batch64 - LR0.01 - MediumDataset - Confusion Matrix

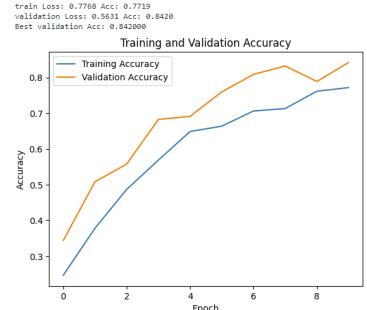
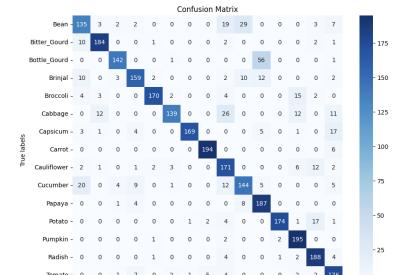


Figure 34. AlexNet - WithTransferLearning - Batch32 - LR0.001 - LargeDataset - Train vs Validation accuracy

	precision	recall	f1-score	support
Bean	0.73	0.68	0.70	200
Bitter_Gourd	0.90	0.92	0.91	200
Bottle_Gourd	0.93	0.71	0.80	200
Brinjal	0.85	0.80	0.82	200
Broccoli	0.96	0.85	0.90	200
Cabbage	0.94	0.69	0.80	200
Capsicum	0.99	0.84	0.91	200
Carrot	0.97	0.97	0.97	200
Cauliflower	0.68	0.85	0.76	200
Cucumber	0.75	0.72	0.74	200
Papaya	0.71	0.94	0.80	200
Potato	0.98	0.87	0.92	200
Pumpkin	0.83	0.97	0.90	200
Radish	0.83	0.94	0.88	200
Tomato	0.76	0.88	0.81	200
accuracy			0.84	3000
macro avg	0.85	0.84	0.84	3000
weighted avg	0.85	0.84	0.84	3000

Figure 35. AlexNet - WithTransferLearning - Batch32 - LR0.001 - LargeDataset - Metrics



## COMP 6721 Winter 2024 Submission #Group I. CONFIDENTIAL REVIEW COPY. DO NOT DISTRIBUTE.

```

1080
1081
1082
1083
1084
1085
1086
1087
1088
1089
1090
1091
1092
1093
1094
1095
1096
1097
1098
1099
1100
1101
1102
1103
1104
1105
1106
1107
1108
1109
1110
1111
1112
1113
1114
1115
1116
1117
1118
1119
1120
1121
1122
1123
1124
1125
1126
1127
1128
1129
1130
1131
1132
1133

```

Train vs Validation Accuracies

```

from sklearn.metrics import classification_report
print(classification_report(y_true, y_pred, target_names=['freshbanana', 'freshokra'], digits=4))

precision    recall   f1-score   support
freshbanana  0.0000  0.0000  0.0000    281
freshokra    0.1517  0.0000  0.0000     67
accuracy      0.0000  0.0000  0.0000    348
macro avg     0.0000  0.0000  0.0000    348
weighted avg  0.0000  0.0000  0.0000    348

```

accuracy  
macro avg  
weighted avg

Figure 37. AlexNet - NoTransferLearning - Batch64 - LR0.01 - Small - Train vs Validation accuracy

```

1114
1115
1116
1117
1118
1119
1120
1121
1122
1123
1124
1125
1126
1127
1128
1129
1130
1131
1132
1133

```

Train vs Validation Accuracies

```

from sklearn.metrics import classification_report
print(classification_report(y_true, y_pred, target_names=['freshbanana', 'freshokra'], digits=4))

precision    recall   f1-score   support
freshbanana  0.0000  0.0000  0.0000    281
freshokra    0.0000  0.0000  0.0000     67
accuracy      0.0000  0.0000  0.0000    348
macro avg     0.0000  0.0000  0.0000    348
weighted avg  0.0000  0.0000  0.0000    348

```

accuracy  
macro avg  
weighted avg

Figure 38. VGGNet - NoTransferLearning - Batch64 - LR0.01 - Small - Train vs Validation accuracy

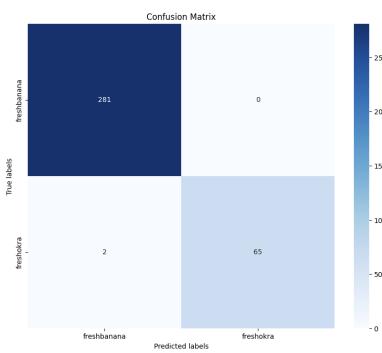


Figure 41. ResNet - NoTransferLearning - Batch64 - LR0.01 - Small - Confusion Matrix

```

1120
1121
1122
1123
1124
1125
1126
1127
1128
1129
1130
1131
1132
1133

```

Train vs Validation Accuracies

Figure 39. ResNet - NoTransferLearning - Batch64 - LR0.01 - Small - Train vs Validation accuracy

```

1134
1135
1136
1137
1138
1139
1140
1141
1142
1143
1144
1145
1146
1147
1148
1149
1150
1151
1152
1153
1154
1155
1156
1157
1158
1159
1160
1161
1162
1163
1164
1165
1166
1167
1168
1169
1170
1171
1172
1173
1174
1175
1176
1177
1178
1179
1180
1181
1182
1183
1184
1185
1186
1187

```

	precision	recall	f1-score	support
freshbanana	0.9929	1.0000	0.9965	281
freshokra	1.0000	0.9701	0.9848	67
accuracy			0.9943	348
macro avg	0.9965	0.9851	0.9907	348
weighted avg	0.9943	0.9943	0.9942	348

Figure 40. ResNet - NoTransferLearning - Batch64 - LR0.01 - Small - Metrics